














MySQL

Let's begin at 19:05 Minsk time

MySQL is an open-source relational database management system (RDBMS)





















- **SQL - structured query language**
- **Data integrity**
- **Fault tolerance**






















341 systems in ranking, December 2018

Dec 2018	Rank		DBMS	Database Model	Score		
	Nov 2018	Dec 2017			Dec 2018	Nov 2018	Dec 2017
1.	1.	1.	Oracle 	Relational DBMS	1283.22	-17.89	-58.32
2.	2.	2.	MySQL 	Relational DBMS	1161.25	+1.36	-156.82
3.	3.	3.	Microsoft SQL Server 	Relational DBMS	1040.34	-11.21	-132.14
4.	4.	4.	PostgreSQL 	Relational DBMS	460.64	+20.39	+75.21
5.	5.	5.	MongoDB 	Document store	378.62	+9.14	+47.85
6.	6.	6.	IBM Db2 	Relational DBMS	180.75	+0.87	-8.83
7.	7.	 8.	Redis 	Key-value store	146.83	+2.66	+23.59
8.	8.	 10.	Elasticsearch 	Search engine	144.70	+1.24	+24.92
9.	9.	 7.	Microsoft Access	Relational DBMS	139.51	+1.08	+13.63
10.	10.	 11.	SQLite 	Relational DBMS	123.02	+0.31	+7.82

Oct 2020	Rank		DBMS	Database Model	Score		
	Sep 2020	Oct 2019			Oct 2020	Sep 2020	Oct 2019
1.	1.	1.	Oracle	Relational, Multi-model	1368.77	-0.59	+12.89
2.	2.	2.	MySQL	Relational, Multi-model	1256.38	-7.87	-26.69
3.	3.	3.	Microsoft SQL Server	Relational, Multi-model	1043.12	-19.64	-51.60
4.	4.	4.	PostgreSQL	Relational, Multi-model	542.40	+0.12	+58.49
5.	5.	5.	MongoDB	Document, Multi-model	448.02	+1.54	+35.93
6.	6.	6.	IBM Db2	Relational, Multi-model	161.90	+0.66	-8.87
7.	8.	7.	Elasticsearch	Search engine, Multi-model	153.84	+3.35	+3.67
8.	7.	8.	Redis	Key-value, Multi-model	153.28	+1.43	+10.37
9.	9.	11.	SQLite	Relational	125.43	-1.25	+2.80
10.	10.	10.	Cassandra	Wide column	119.10	-0.08	-4.12

Top 10 DBMS systems in ranking, February 2022

Rank			DBMS	Database Model	Score		
Feb 2022	Jan 2022	Feb 2021			Feb 2022	Jan 2022	Feb 2021
1.	1.	1.	Oracle 	Relational, Multi-model 	1256.83	-10.05	-59.84
2.	2.	2.	MySQL 	Relational, Multi-model 	1214.68	+8.63	-28.69
3.	3.	3.	Microsoft SQL Server 	Relational, Multi-model 	949.05	+4.24	-73.88
4.	4.	4.	PostgreSQL  	Relational, Multi-model 	609.38	+2.83	+58.42
5.	5.	5.	MongoDB 	Document, Multi-model 	488.64	+0.07	+29.69
6.	6.	 7.	Redis 	Key-value, Multi-model 	175.80	-2.18	+23.23
7.	7.	 6.	IBM Db2	Relational, Multi-model 	162.88	-1.32	+5.26
8.	8.	8.	Elasticsearch	Search engine, Multi-model 	162.29	+1.54	+11.29
9.	9.	 11.	Microsoft Access	Relational	131.26	+2.31	+17.09
10.	10.	 9.	SQLite 	Relational	128.37	+0.94	+5.20

Rank			DBMS	Database Model	Score		
Jan 2025	Dec 2024	Jan 2024			Jan 2025	Dec 2024	Jan 2024
1.	1.	1.	Oracle 	Relational, Multi-model 	1258.76	-5.03	+11.27
2.	2.	2.	MySQL 	Relational, Multi-model 	998.15	-5.61	-125.31
3.	3.	3.	Microsoft SQL Server	Relational, Multi-model 	798.55	-7.14	-78.05
4.	4.	4.	PostgreSQL 	Relational, Multi-model 	663.41	-2.97	+14.45
5.	5.	5.	MongoDB 	Document, Multi-model 	402.50	+2.12	-14.98
6.	 7.	 9.	Snowflake 	Relational	153.90	+6.54	+27.98
7.	 6.	 6.	Redis 	Key-value, Multi-model 	153.36	+3.08	-6.03
8.	8.	 7.	Elasticsearch	Multi-model 	134.92	+2.60	-1.15
9.	9.	 8.	IBM Db2	Relational, Multi-model 	122.97	+0.19	-9.43
10.	10.	 11.	SQLite	Relational	106.69	+4.97	-8.51

Method of calculating the scores of the DB-Engines Ranking

The DB-Engines Ranking is a list of database management systems ranked by their current popularity. We measure the popularity of a system by using the following parameters:

- **Number of mentions of the system on websites**, measured as number of results in search engines queries. At the moment, we use [Google](#) and [Bing](#) for this measurement. In order to count only relevant results, we are searching for <system name> together with the term database, e.g. "Oracle" and "database".
- **General interest in the system**. For this measurement, we use the frequency of searches in [Google Trends](#).
- **Frequency of technical discussions about the system**. We use the number of related questions and the number of interested users on the well-known IT-related Q&A sites [Stack Overflow](#) and [DBA Stack Exchange](#).
- **Number of job offers, in which the system is mentioned**. We use the number of offers on the leading job search engines [Indeed](#) and [Simply Hired](#).
- **Number of profiles in professional networks, in which the system is mentioned**. We use the internationally most popular professional network [LinkedIn](#).
- **Relevance in social networks**. We count the number of [Twitter](#) tweets, in which the system is mentioned.

MySQL: GitHub, US Navy, NASA, Tesla, Netflix, WeChat, Facebook, Zendesk, Twitter, Zappos, YouTube, Spotify.

Oracle: Bauerfeind AG, CAIRN India, Capcom Co., ChevronTexaco, Coca-Cola FEMSA, COOP Switzerland, ENEL, Heidelberger Druck, MTU Aero Engines, National Foods Australia, Spire Healthcare, Stadtwerke München, Swarovski, Tyson Foods, TVS Motor Company, Vilene.

- Data types;
- Architecture;
- Indexes;
- ACID

Numeric Data Types

- **TINYINT**: from -127 to 128, takes 1 byte
- **BOOL (BOOLEAN)**: `TINYINT(1)`.
- **TINYINT UNSIGNED**: from 0 to 255, takes 1 byte
- **SMALLINT**: from -32768 to 32767, takes 2 bytes
- **SMALLINT UNSIGNED**: from 0 to 65535, takes 2 bytes
- **MEDIUMINT**: from -8388608 to 8388607, takes 3 bytes
- **MEDIUMINT UNSIGNED**: from 0 to 16777215, takes 3 bytes
- **INT**: from -2147483648 to 2147483647, takes 4 bytes
- **INT UNSIGNED**: from 0 to 4294967295, takes 4 bytes
- **BIGINT**: from -9 223 372 036 854 775 808 to 9 223 372 036 854 775 807, takes 8 bytes
- **BIGINT UNSIGNED**: from 0 to 18 446 744 073 709 551 615, takes 8 bytes

Numeric Data Types

- **DECIMAL**: Fixed-precision number. `DECIMAL(precision, scale)`. $1 < \text{precision} < 65$
- **FLOAT**: $-3.4028 * 10^{38}$ to $3.4028 * 10^{38}$, takes 4 bytes `FLOAT(M,D)`
- **DOUBLE**: from $-1.7976 * 10^{308}$ to $1.7976 * 10^{308}$, takes 8 bytes. `DOUBLE(M,D)`

Data types for a work with date and time

- **DATE**: stores dates from JAN 1 1000 to DEC 31 9999. By default it stores data using yyyy-mm-dd format. Takes 3 bytes.
- **TIME**: stores time from -838:59:59 to 838:59:59. By default it stores data using "hh:mm:ss" format. Takes 3 bytes.
- **DATETIME**: from "1000-01-01 00:00:00" to "9999-12-31 23:59:59". By default it stores data using "yyyy-mm-dd hh:mm:ss" format. Takes 8 bytes
- **TIMESTAMP**: from "1970-01-01 00:00:01" UTC to "2038-01-19 03:14:07" UTC. Takes 4 bytes
- **YEAR**: stores a year as 4 digits. Takes 1 byte.

String Types

Can build index

- **CHAR**: fixed length string
- **VARCHAR**: variable length string

Can build index with limited length

- **TINYTEXT**: text up to 255 bytes.
- **TEXT**: text up to 65 KB.
- **MEDIUMTEXT**: text up to 16 MB
- **LARGETEXT**: text up to GB




Union Types

- **ENUM**: stores single value from the list of allowed values. Takes 1-2 bytes

```
CREATE TABLE tickets (  
    id INT PRIMARY KEY AUTO_INCREMENT,  
    title VARCHAR(255) NOT NULL,  
    priority ENUM('Low', 'Medium', 'High') NOT NULL  
);
```

```
INSERT INTO tickets(title, priority) VALUES('Scan virus for computer A', 'High');
```

```
INSERT INTO tickets(title, priority) VALUES('Upgrade Windows OS for all computers', 1);
```




	 id	 title	 priority
1	1	Scan virus for computer A	High
2	2	Upgrade Windows OS for all computers	Low

Union Types

- **SET**: stores multiple values from the list of allowed values (up to 64 possible values). Takes 1-8 bytes.

```
CREATE TABLE tickets2 (  
    id INT PRIMARY KEY AUTO_INCREMENT,  
    title VARCHAR(255) NOT NULL,  
    priority SET('2^0', '2^1', '2^2') NOT NULL  
);
```

```
INSERT INTO tickets2(title, priority) VALUES('Show example', '2^0,2^1');  
INSERT INTO tickets2(title, priority) VALUES('Show example 2', 5);
```

	 id	 title	 priority
1	1	Show example	2^0,2^1
2	2	Show example 2	2^0,2^2

Binary data types

- **TINYBLOB**: binary text up to 255 bytes.
- **BLOB**: binary text up to 65 KB.
- **MEDIUMBLOB**: binary text up to 16 MB
- **LARGEBLOB**: binary text up to 4 GB

JSON data type

Native JSON support which includes

1. Automatic JSON validation
2. Optimized storage
3. Partial update using
 1. JSON_SET
 2. JSON_INSERT
 3. JSON_REPLACE
 4. ...

ORM

```
// Create a new user  
const jane = await User.create({ firstName: "Jane", lastName: "Doe" });  
console.log("Jane's auto-generated ID:", jane.id);
```

```
const { Op } = require("sequelize");  
Post.findAll({  
  where: {  
    authorId: {  
      [Op.eq]: 2  
    }  
  }  
});  
// SELECT * FROM post WHERE authorId = 2;
```

```
// Change everyone without a last name to "Doe"  
await User.update({ lastName: "Doe" }, {  
  where: {  
    lastName: null  
  }  
});
```

CREATE TABLE Users

(
 id INT UNSIGNED NOT NULL PRIMARY KEY AUTO_INCREMENT,
 user_id INT NOT NULL,
 user_name CHAR(64) NOT NULL
);

DROP TABLE Users;

ALTER TABLE Users

MODIFY COLUMN user_name VARCHAR(64) NOT NULL,
 ADD COLUMN group_id TINYINT(3) DEFAULT NULL AFTER user_id,
 ADD INDEX name (user_name),
 DROP COLUMN surname,
 ADD CONSTRAINT FK_Group FOREIGN KEY (group_id) REFERENCES Groups (group_id);

```
SELECT * FROM Users (WHERE user_id = 1);
```

```
UPDATE Users SET user_name = 'Иванов И.И.' WHERE user_id = 1;
```

```
INSERT INTO Users (user_id, group_id, user_name) VALUES (2, 1, 'Петров П.П.');
```


```
INSERT INTO Users VALUES (3, 1, 'Петрович П.П. '), (4, NULL, 'Сидоров С.С.');
```

```
DELETE FROM Users WHERE user_id = 1;
```


```
SELECT * FROM Users a
  INNER JOIN Groups g
    ON a.group_id = g.group_id
 WHERE user_id = 1;
```

```
SELECT * FROM Users a
  LEFT JOIN Groups g
    ON a.group_id = g.group_id
 WHERE user_id = 1;
```


```
SELECT <fields>
FROM TableA A
INNER JOIN TableB B
ON A.key = B.key
```



```
SELECT <fields>
FROM TableA A
RIGHT JOIN TableB B
ON A.key = B.key
```




```
SELECT <fields>
FROM TableA A
LEFT JOIN TableB B
ON A.key = B.key
```




SQL

JOINS


```
SELECT <fields>
FROM TableA A
RIGHT JOIN TableB B
ON A.key = B.key
WHERE a.key IS NULL
```




```
SELECT <fields>
FROM TableA A
LEFT JOIN TableB B
ON A.key = B.key
WHERE B.key IS NULL
```



```
SELECT <fields>
FROM TableA A
FULL OUTER JOIN TableB B
ON A.key = B.key
```



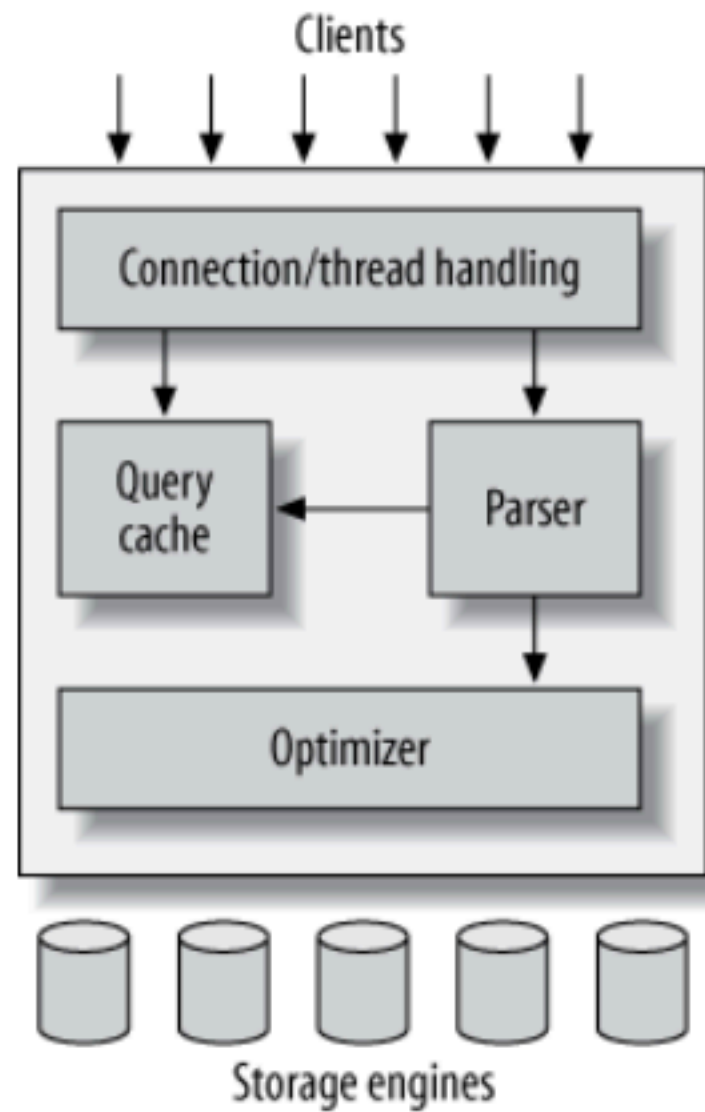
```
SELECT <fields>
FROM TableA A
FULL OUTER JOIN TableB B
ON A.key = B.key
WHERE A.key IS NULL
OR B.key IS NULL
```



List groups which have more than 2 users excluding test users

```
SELECT g.group_name, COUNT(*) as cnt
FROM Users u
INNER JOIN Groups g
    ON u.group_id = g.id
WHERE u.is_test=0
GROUP BY g.id
HAVING cnt > 2;
```

Architecture



Storage engines


- **MyISAM**
- **Memory**
- **CSV**
- **Archive**
- **Blackhole**

InnoDB



- **FOREIGN KEY**
- **Transactions**
- **ACID**
 - **Atomicity**
 - **Consistency**
 - **Isolation**
 - **Durability**
- **Buffer pool**
- **Change buffer**
- **Double write buffer**

- **Primary Key**
- **Unique**
- **Index**

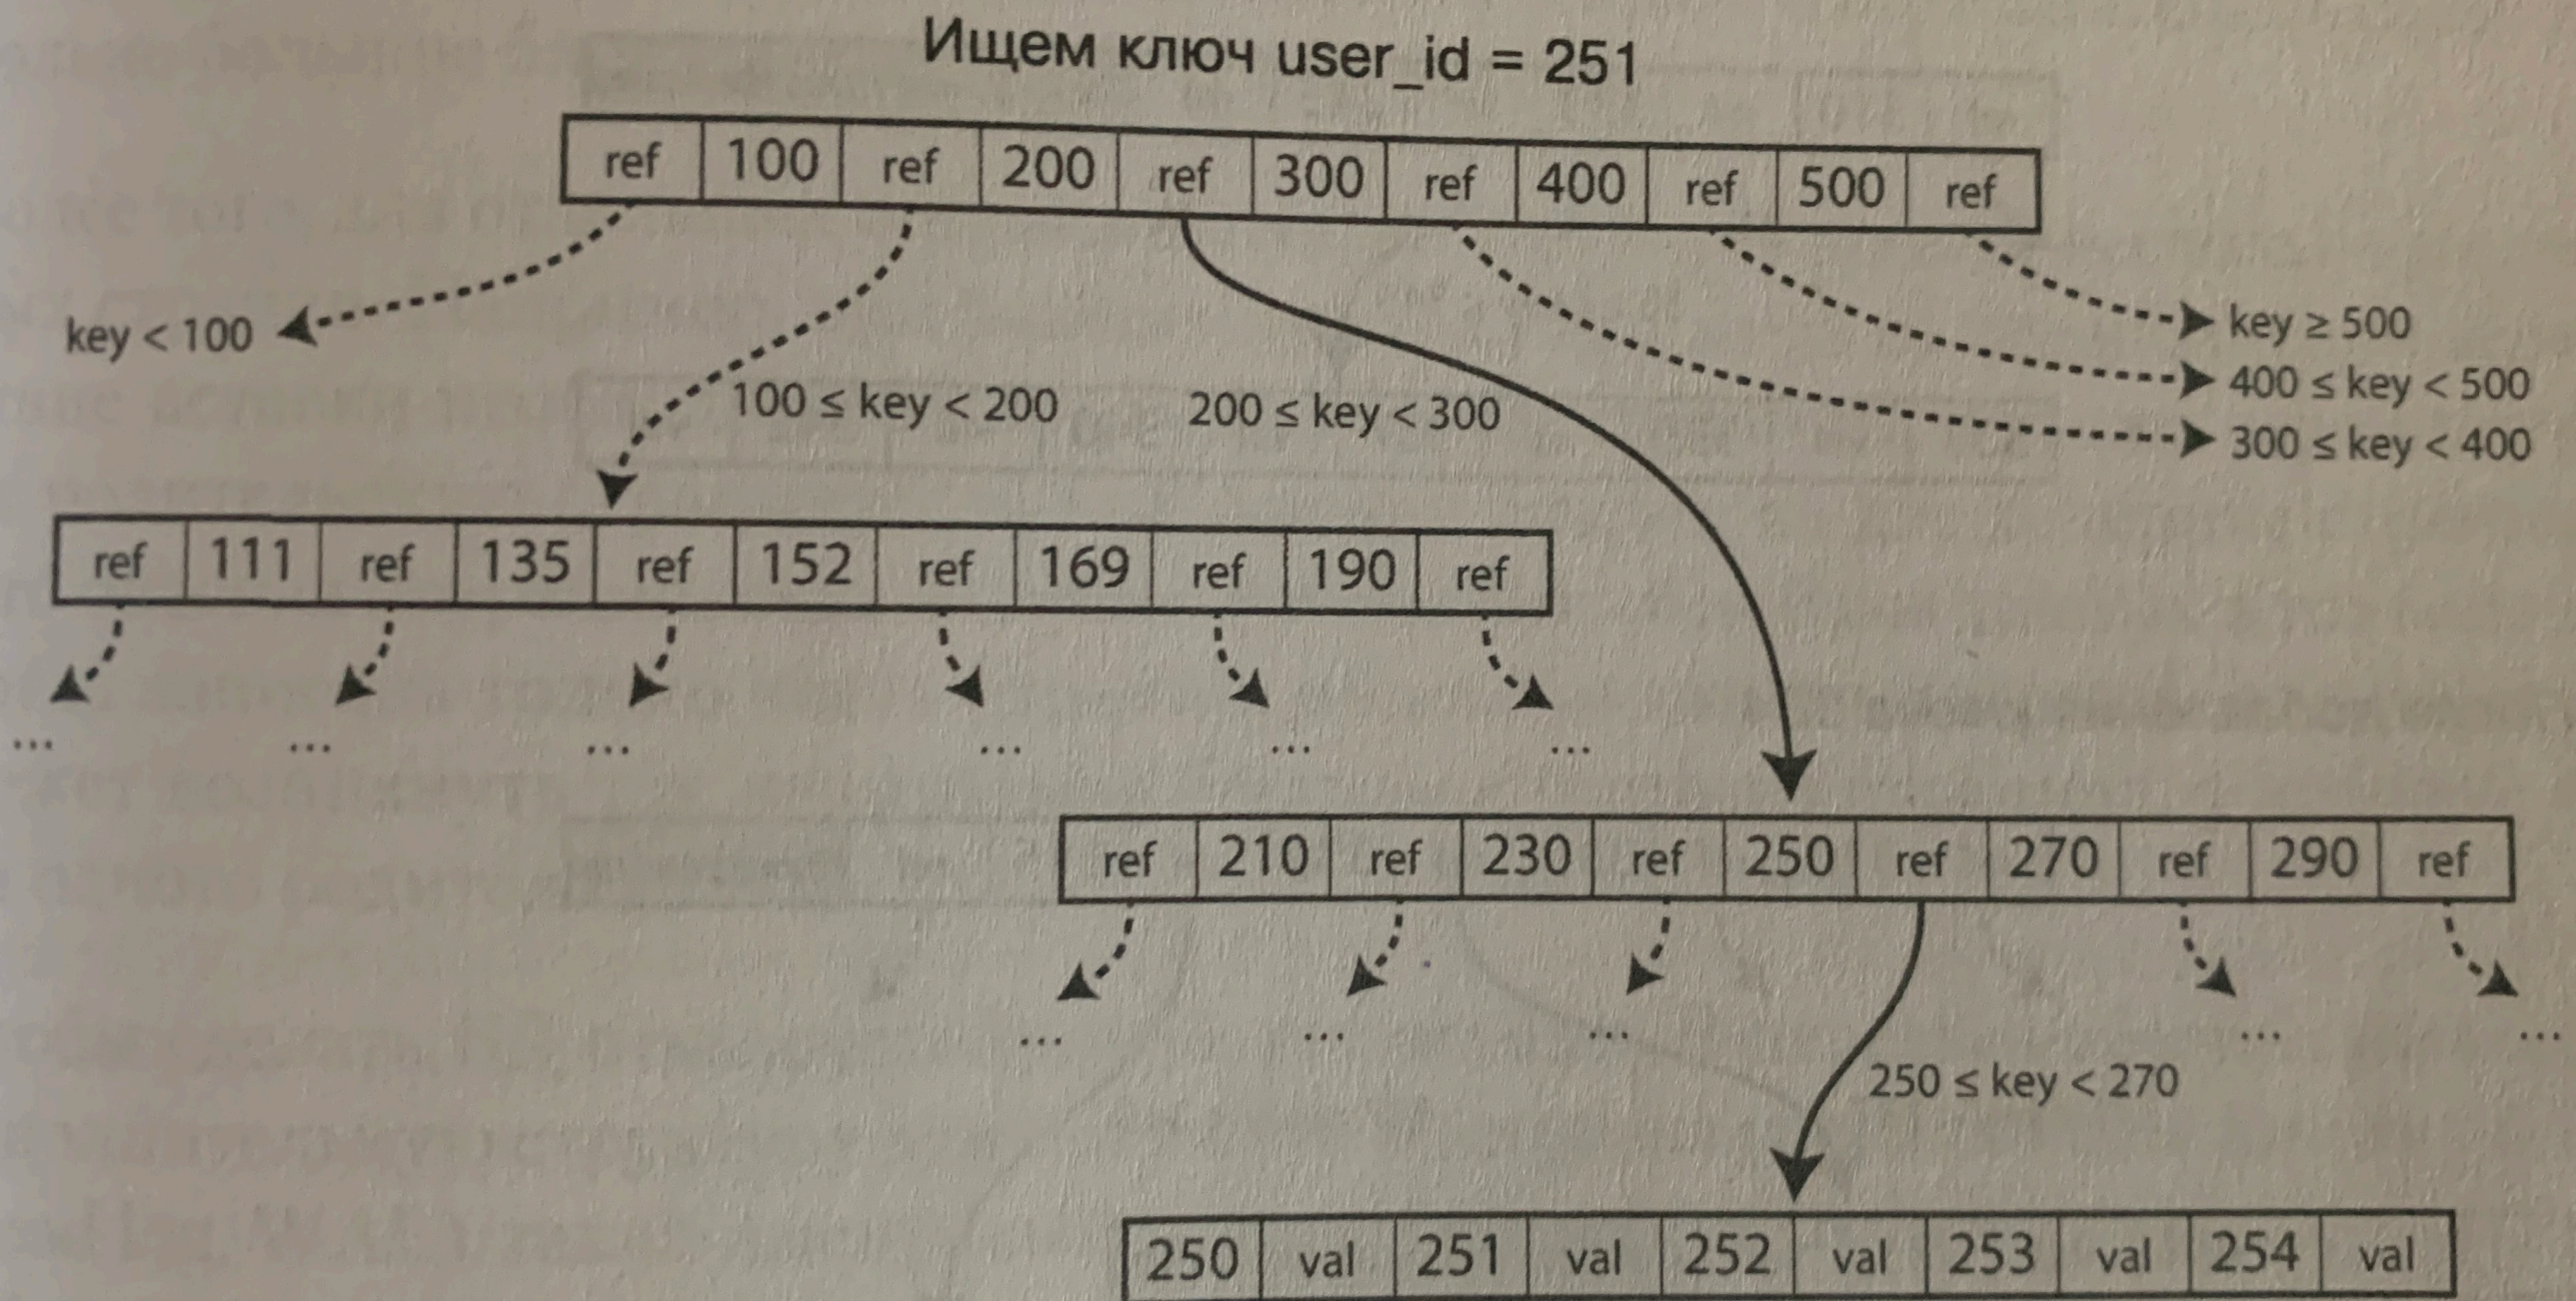
B-tree Index



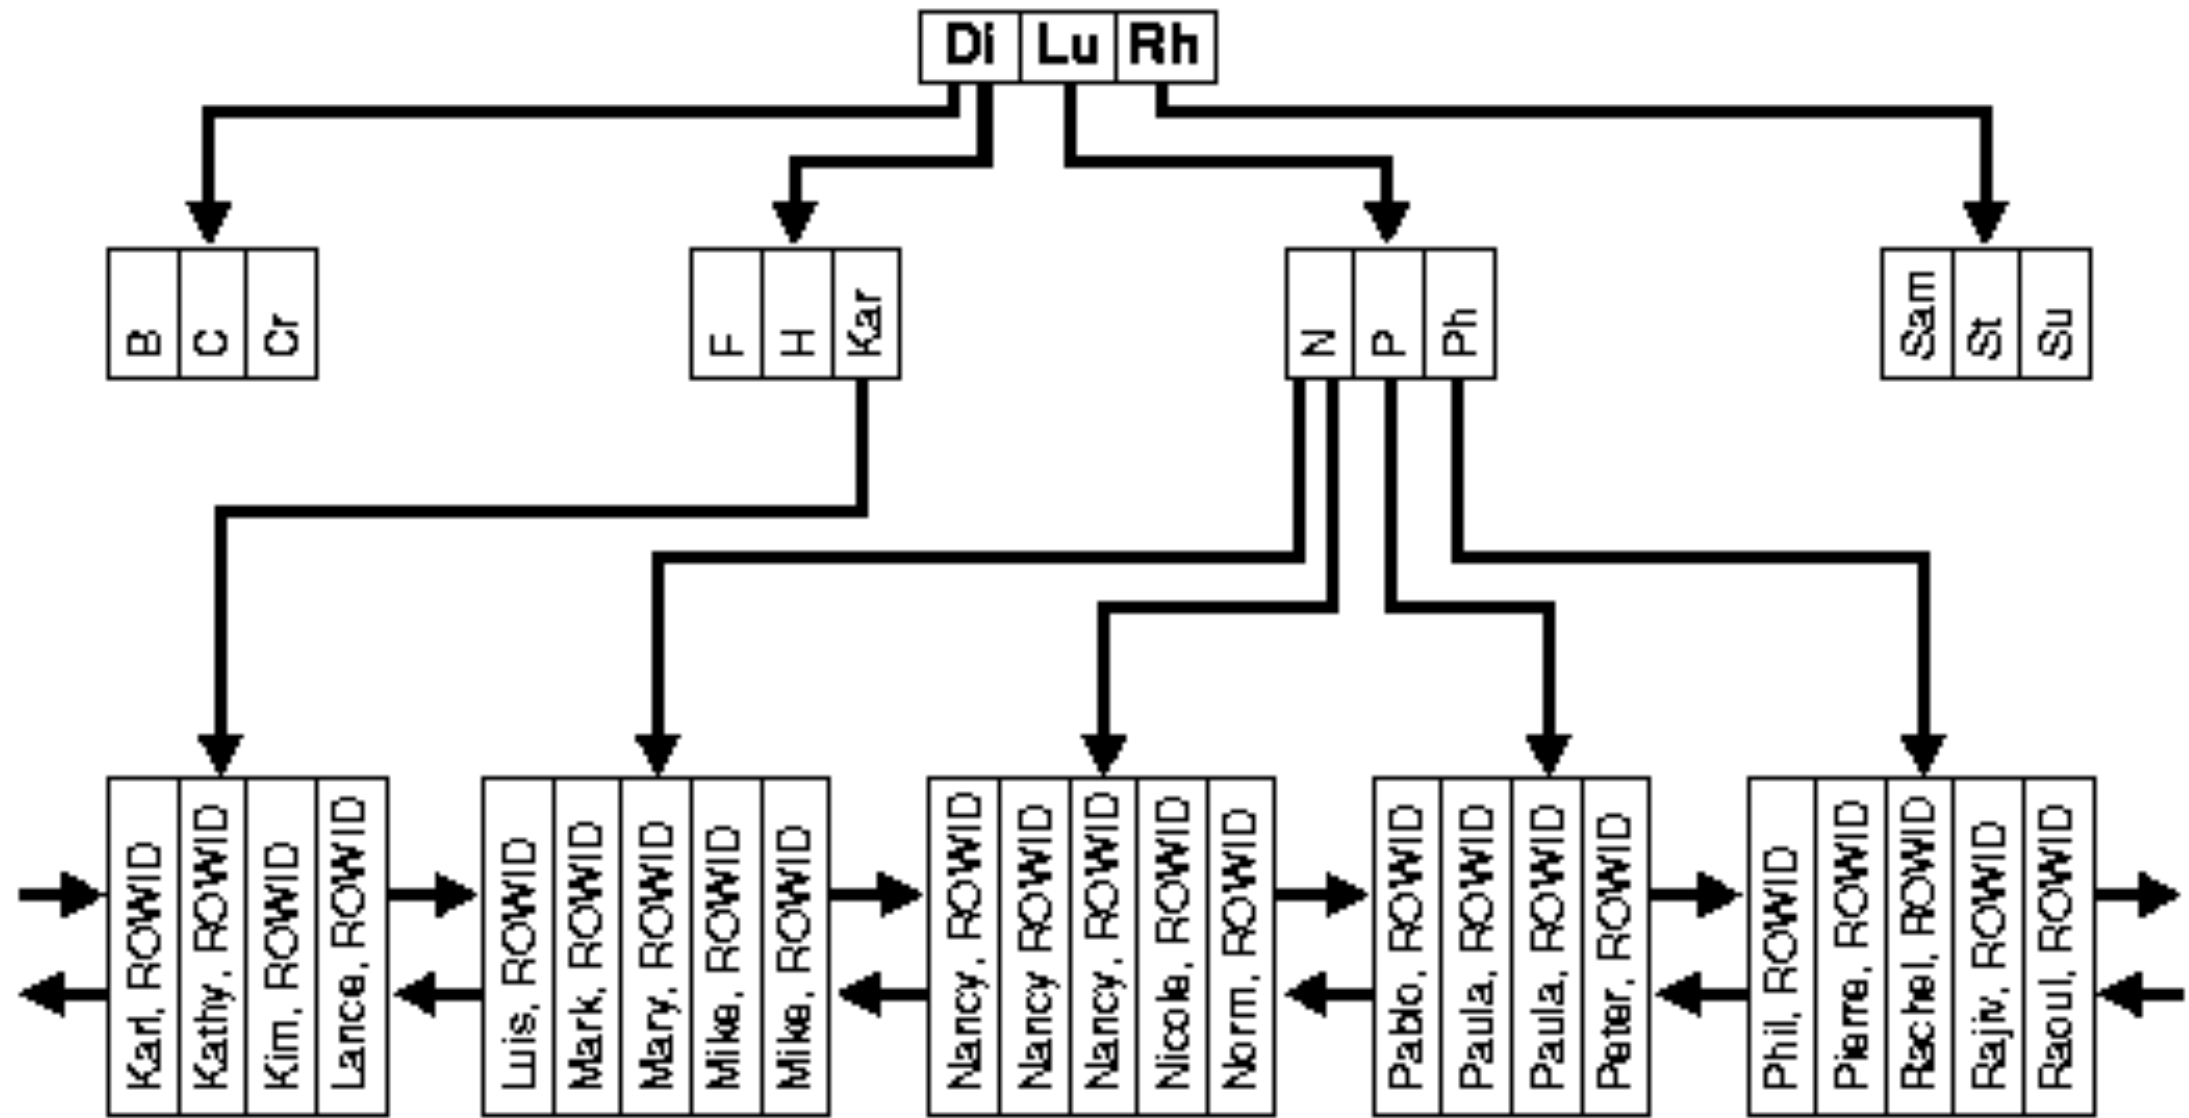
B-tree Index



B-tree Index



B-tree Index



Complex index

```
SELECT * FROM users WHERE age = 29 AND gender = 'male'
```

```
CREATE INDEX age_gender ON users(age, gender);
```

```
SELECT * FROM users WHERE age >= 10 AND gender = 'male'
```

```
CREATE INDEX age_gender ON users(gender, age);
```


Selectivity

The selectivity basically is a measure of how much variety there is in the values of a given table column in relation to the total number of rows in a given table

InnoDB

- **FOREIGN KEY**
- **Transactions**
- **ACID**
 - **Atomicity**
 - **Consistency**
 - **Isolation**
 - **Durability**
- **Buffer pool**
- **Change buffer**
- **Double write buffer**

Isolation levels

Initial state

```
CREATE TABLE accounts (  
  id INT PRIMARY KEY,  
  name VARCHAR(50),  
  balance INT  
);  
INSERT INTO accounts (id, name, balance)  
VALUES (id 1, name 'Alice', balance 100);
```

Isolation levels

READ UNCOMMITTED

A transaction can see changes to data made by other transactions that are not committed yet. This can lead to **dirty reads**.

Transaction A

```
-- Transaction A (any isolation level)
BEGIN;
UPDATE accounts SET balance = 200 WHERE id = 1; -- No commit yet
-- Some other operations
COMMIT;
```

Transaction B

```
-- Transaction B
SET TRANSACTION ISOLATION LEVEL READ UNCOMMITTED;
BEGIN;

-- First Read (before Transaction A starts)
SELECT balance FROM accounts WHERE id = 1; -- Alice's balance
-- Output: 100

-- Second Read (Transaction A has updated but not committed)
SELECT balance FROM accounts WHERE id = 1;
-- Output: 200 (reads uncommitted update)

-- Third Read (Transaction A has committed)
SELECT balance FROM accounts WHERE id = 1;
-- Output: 200
```

Isolation levels

READ COMMITTED

Only committed changes are visible. Prevents **dirty reads**.

Transaction A

```
-- Transaction A (any isolation level)
BEGIN;
UPDATE accounts SET balance = 200 WHERE id = 1; -- No commit yet
-- Some other operations
COMMIT;
```

Transaction B

```
-- Transaction B
SET TRANSACTION ISOLATION LEVEL READ COMMITTED;
BEGIN;

-- First Read (before Transaction A starts)
SELECT balance FROM accounts WHERE id = 1;
-- Output: 100

-- Second Read (Transaction A has updated but not committed)
SELECT balance FROM accounts WHERE id = 1;
-- Output: 100 (does not see uncommitted update)

-- Third Read (Transaction A has committed)
SELECT balance FROM accounts WHERE id = 1;
-- Output: 200 (reads committed update)
```

Isolation levels

REPEATABLE READ (DEFAULT VALUE)

Each transaction has a consistent snapshot of data, so repeated reads yield the same result. Prevents dirty and non-repeatable reads.

Transaction A

```
-- Transaction A (any isolation level)
BEGIN;
UPDATE accounts SET balance = 200 WHERE id = 1; -- No commit yet
-- Some other operations
COMMIT;
```

Transaction B

```
-- Transaction B
SET TRANSACTION ISOLATION LEVEL REPEATABLE READ;
BEGIN;

-- First Read (before Transaction A starts)
SELECT balance FROM accounts WHERE id = 1;
-- Output: 100

-- Second Read (Transaction A has updated but not committed)
SELECT balance FROM accounts WHERE id = 1;
-- Output: 100 (sees consistent snapshot, ignoring uncommitted changes)

-- Third Read (Transaction A has committed)
SELECT balance FROM accounts WHERE id = 1;
-- Output: 100 (still sees original snapshot)
```

Isolation levels

SERIALIZABLE

The strictest level, each transaction is fully isolated, as if transactions run one at a time.

Transaction A

```
-- Transaction A (any isolation level)
BEGIN;
UPDATE accounts SET balance = 200 WHERE id = 1; -- No commit yet
-- Some other operations
COMMIT;
```

Transaction B

```
-- Transaction B
SET TRANSACTION ISOLATION LEVEL SERIALIZABLE;
BEGIN;

-- First Read (before Transaction A starts)
SELECT balance FROM accounts WHERE id = 1;
-- Output: 100

-- Second Read (Transaction A is trying to update)
SELECT balance FROM accounts WHERE id = 1;
-- Output: 100 (Transaction A is blocked)

-- Third Read (Transaction A is still blocked until Transaction B finish)
SELECT balance FROM accounts WHERE id = 1;
-- Output: 100
```

What we can use

- **Terminal**
- **DataGrip**
- **HeidiSQL**
- **Mysql Workbench**
- **DbViewer**

What to read

- 7 databases in 7 weeks
- High Performance MySQL, 4th Edition
- Designing Data-Intensive Applications - Kleppmann

Questions?